

Precision Time Synchronization in Data Centers

Research Scientist, Meta
Ahmad Byagowi

Why to have precision time synchronization in Data Centers?

How to have precision time synchronization in Data Centers?

Background

- Speed of light and speed of electricity are finite!
 - Speed of Light is slower in fiber ($2.14 \times 10^8 \text{m/s}$) than in vacuum ($2.99 \times 10^8 \text{m/s}$)
 - Latency in data transfer over the network
- Hyperscale cloud services serve people across the globe
 - Geographical distribution of customers
 - Necessity of distribution due to safety, robustness and regulations
- The demand for hardware resources is always increasing
 - Increase in users, data and multi dimensional contents
 - Interactivity with the data like multi-user and VR
 - Big data and AI

Solution to Address the Hyperscale Demand

- Horizontal Scaling
 - Expansion of resources
- Geographical expansion
 - Reduction of latency



Challenges with Expansion

- Distributed Systems
 - Consistency vs Availability
 - Propagation of Information
 - Need for Redundancy
- Distribution of Clock
 - A common reference between all machines
 - Latency vs Clock Skewness

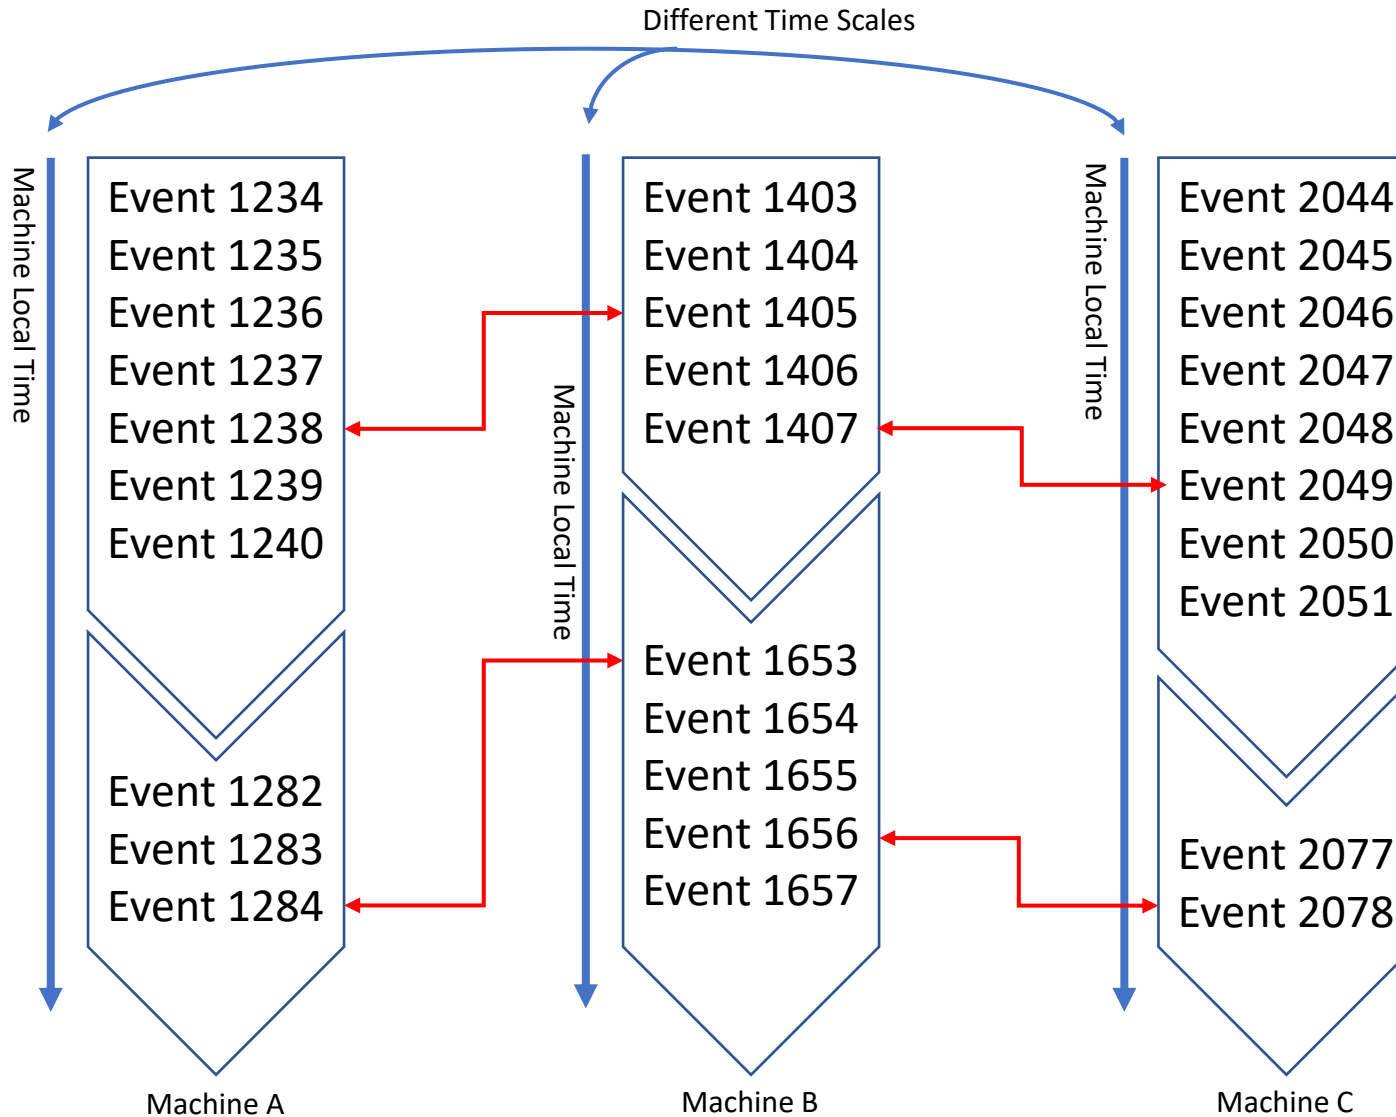
Categories of use cases

Use cases of Precision Time Sync in Distributed Systems:

- Synchronization (Phase)
 - Active (1)
 - Running Events at specific time
 - Sync or desync
 - Reactive (2)
 - Measure latency, time or intervals
- Syntonization (Frequency)
 - Active (3)
 - Calibrate speed and align runtime to reduce tail latency
 - Reactive (4)
 - Measure heterogeneity or provide binning

	Phase	Frequency
Active	1	2
Reactive	3	4

Associate Events Between Multiple Machines



Precision Requirement

- CPU level
- OS (Kernel) level
- Distributed System level

Precision Requirement at CPU level

- Nyquist sampling theorem
 - Sampling interval required to avoid aliasing
 - Sampling frequency should be at least twice the highest frequency contained in the signal
- Frequency in event occurrence
 - Instruction Latency
 - Instruction Throughput

// mov = 1 CPU cycle

// xchg = 3 CPU cycles

// rdtsc = 1 CPU cycle

A CPU with a clock speed of 3.2 GHz executes 3.2 billion cycles per second
That is a period of about 310ps

Precision Requirement at OS level


- dmesg

```
[52603.373642] {38}[Hardware Error]: event severity: corrected
[52603.373643] {38}[Hardware Error]: Error 0, type: corrected
[52603.373644] {38}[Hardware Error]: section_type: PCIe error
[52603.373644] {38}[Hardware Error]: port_type: 4, root port
[52603.373645] {38}[Hardware Error]: version: 3.0
[52603.373645] {38}[Hardware Error]: command: 0x0547, status: 0x0010
[52603.373646] {38}[Hardware Error]: device_id: 0000:b7:01.0
[52603.373647] {38}[Hardware Error]: slot: 255
[52603.373648] {38}[Hardware Error]: secondary_bus: 0xb8
[52603.373648] {38}[Hardware Error]: vendor_id: 0x8086, device_id: 0x352a
[52603.373649] {38}[Hardware Error]: class_code: 060400
[52603.373649] {38}[Hardware Error]: bridge: secondary_status: 0x0000, control: 0x0013
```

System Logging is based on clock_boottime (clock_minotone_RAW) with a quanta on 1us
Events occur faster than the quanta of 1us (aliasing)


Challenges and the Precision Requirement

- Vernier acuity




- Compounding of Events

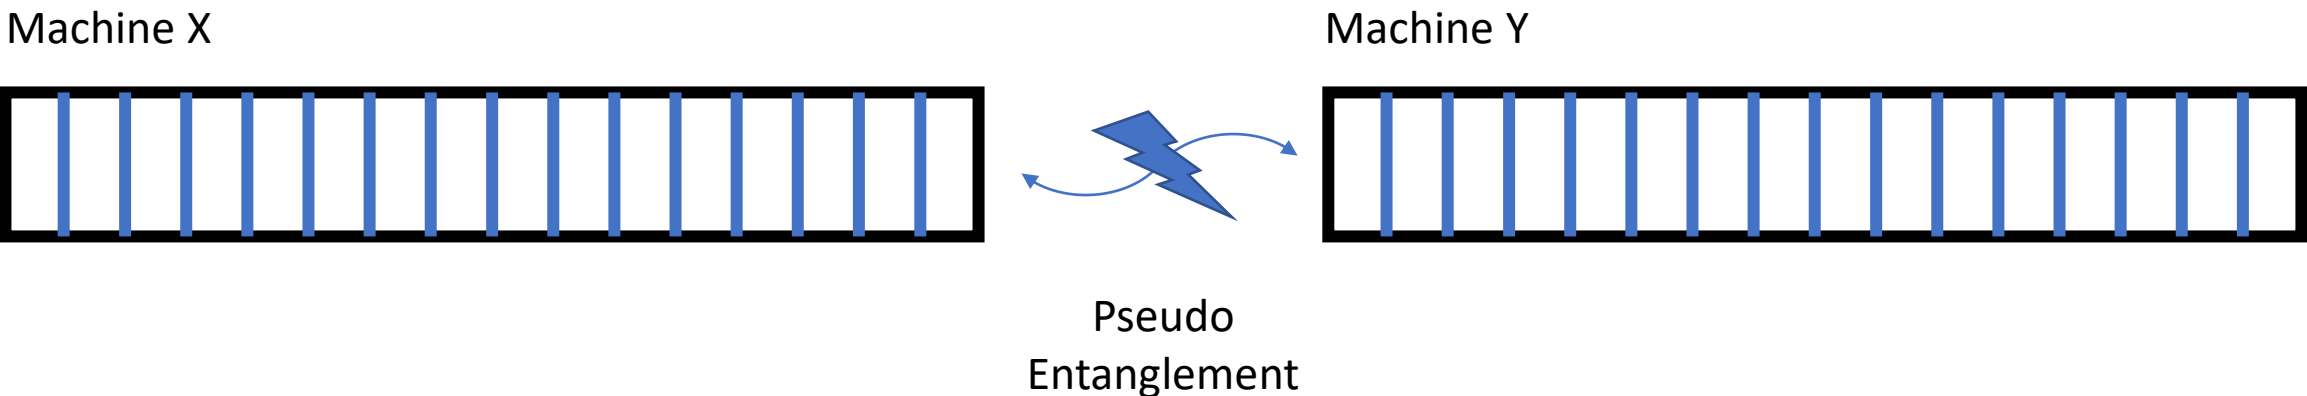
$$1^n = 1$$



Precision Requirement for Distributed Systems




What comes out of Precision Time Sync?



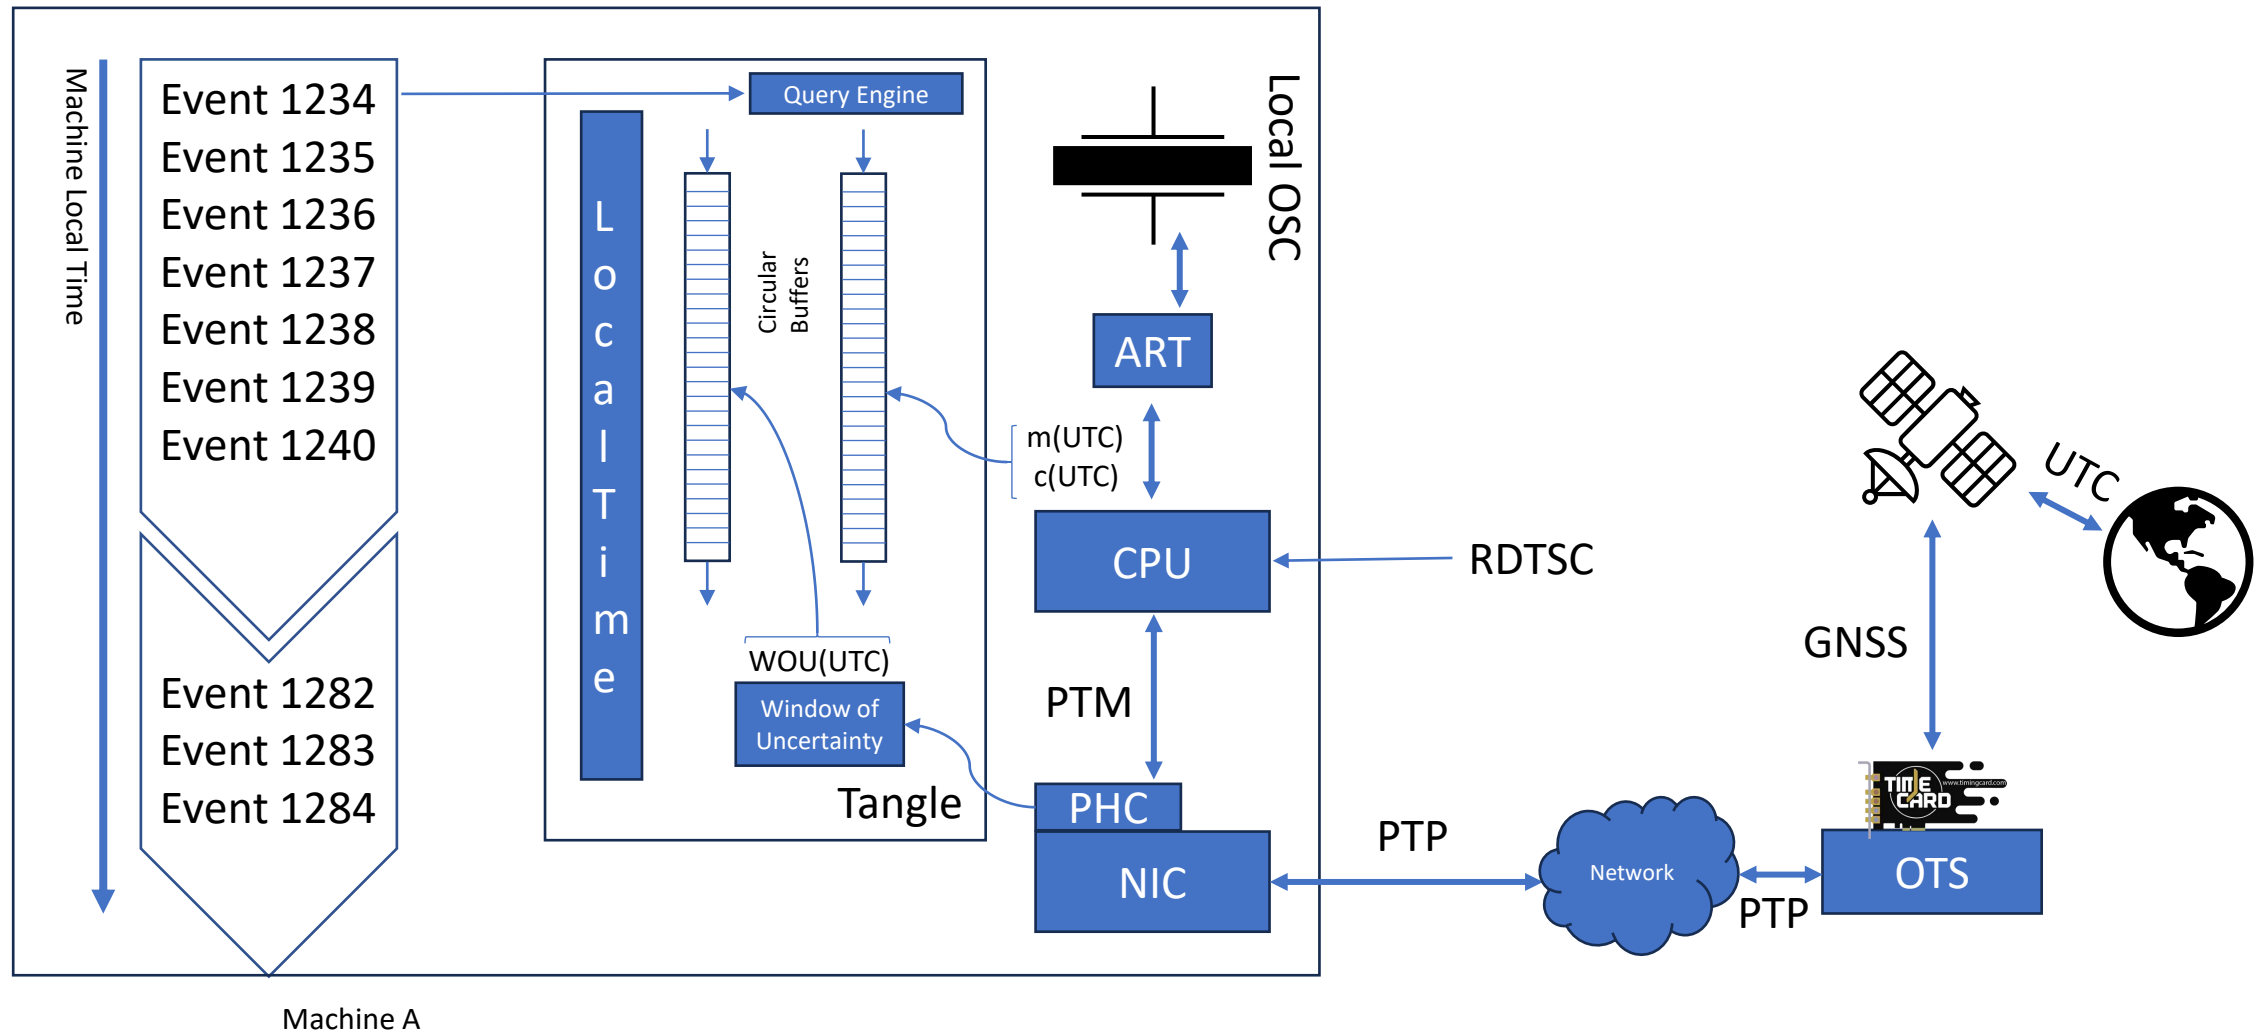
Machine X and Y are any two machines across the globe or inside a local network

Pseudo Entanglement: Probabilistic Entanglement of two Registers (Machine Y and Y) within the Windows of Uncertainty

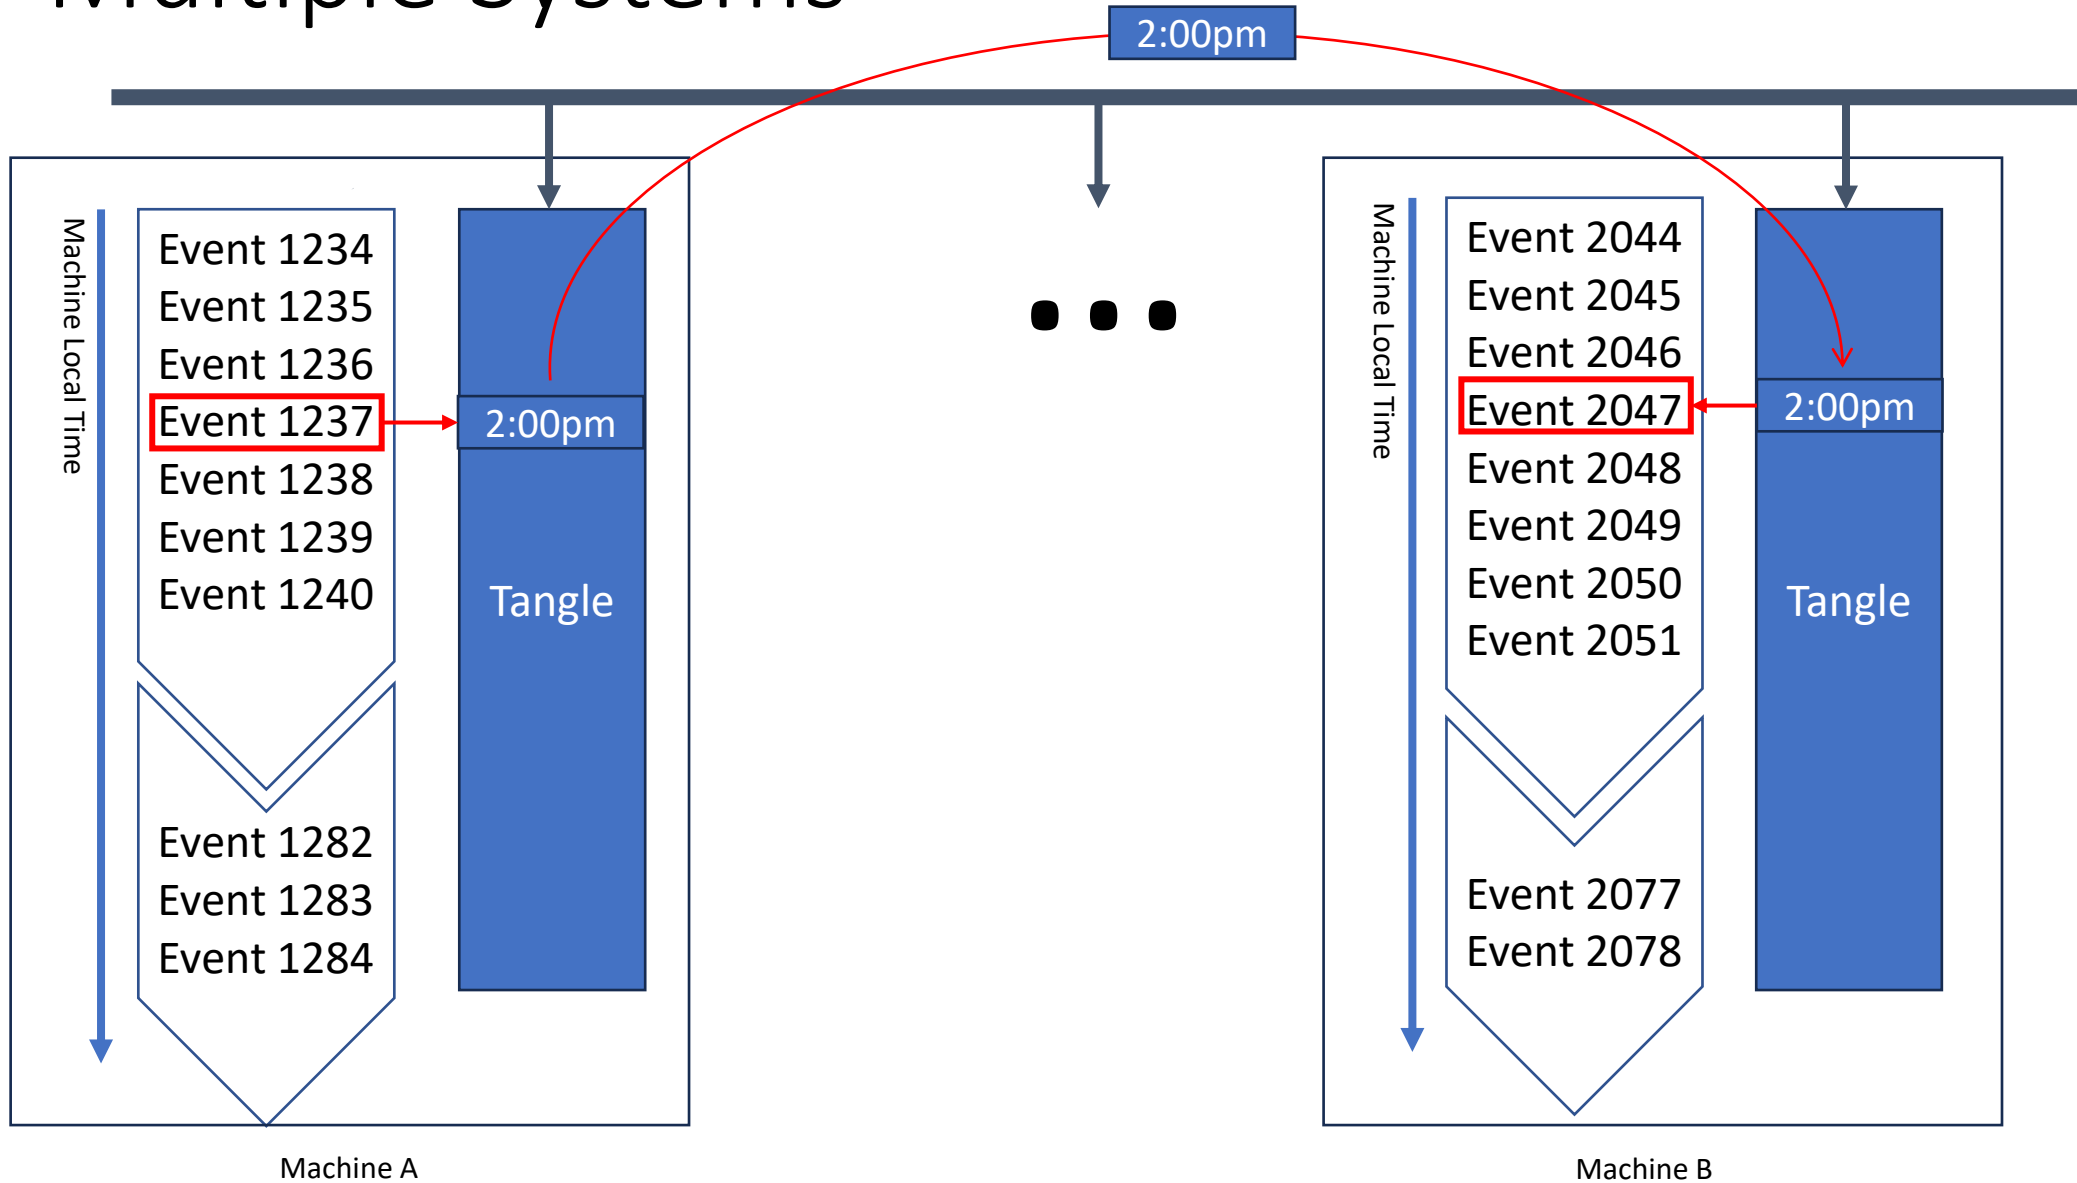
Window of Uncertainty: An ongoing estimation of a time interval that UTC (or TAI) sits inside it (with a given probability)



Tangle



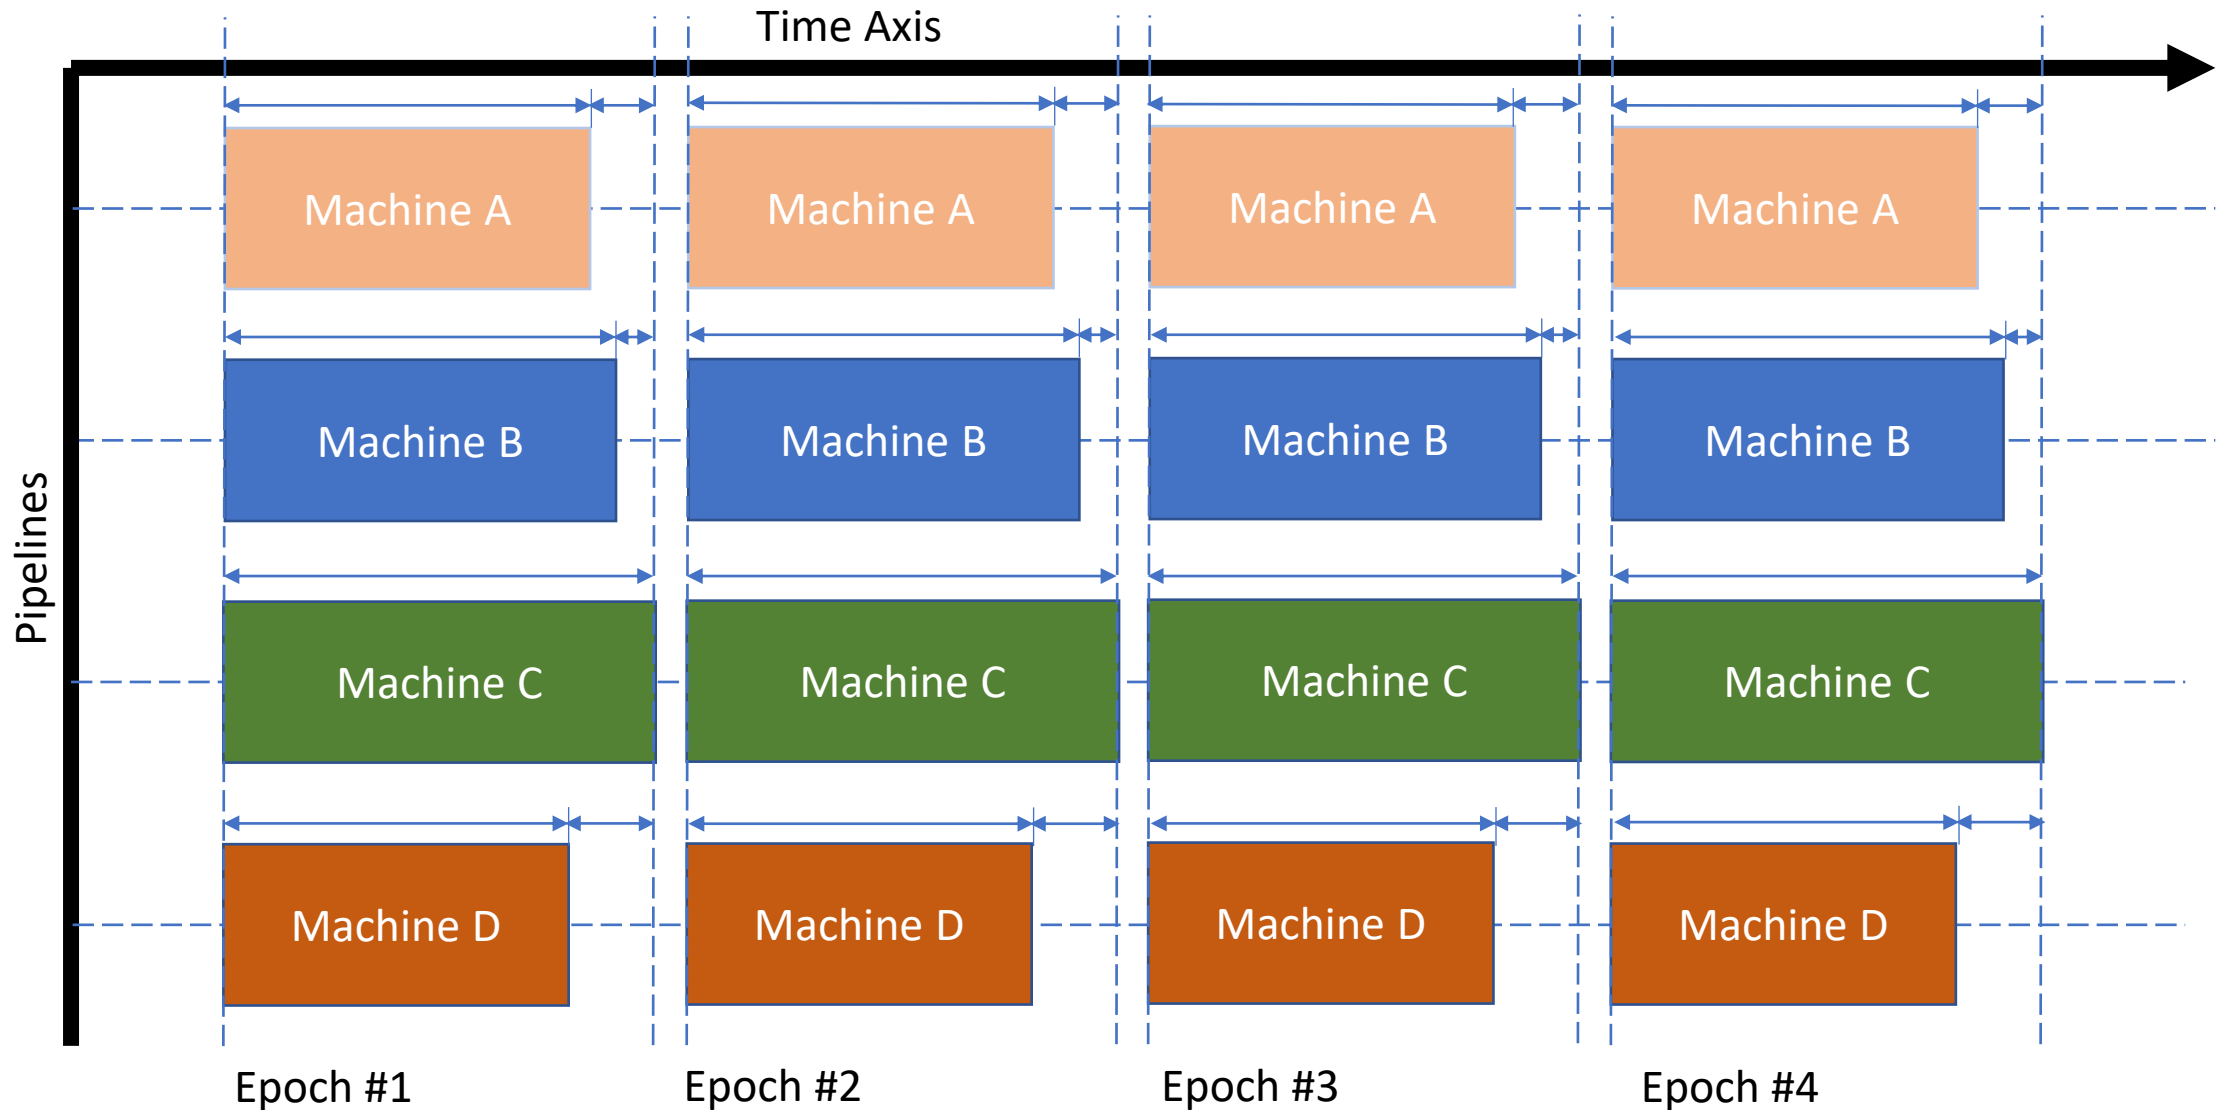
Multiple Systems



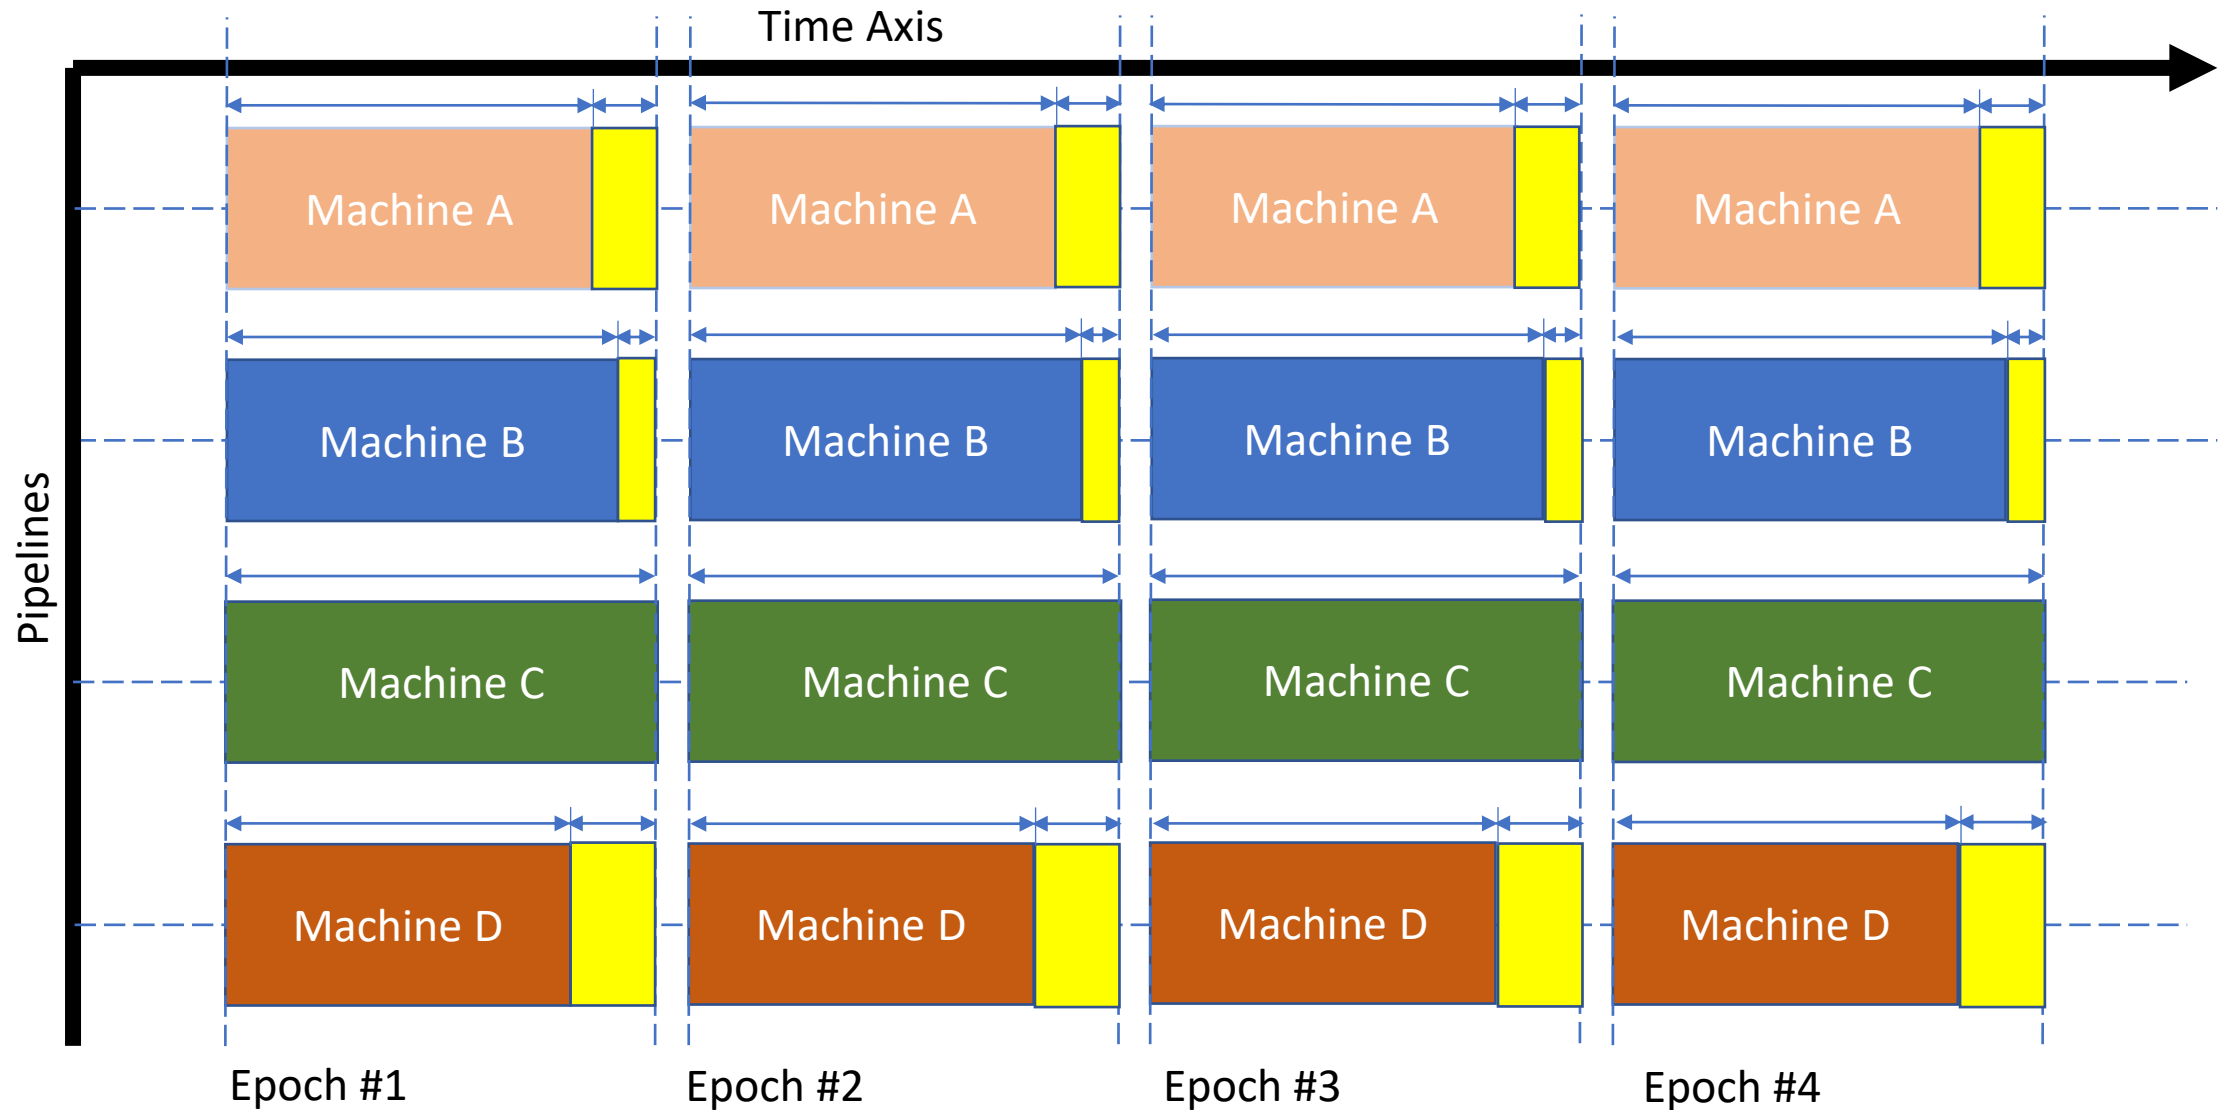
Functions

- Identify concurrent event in another machine[s]
- Find the timestamp of an event in another machine[s]
- Chronologically Rank a given event across machines
- Measure the one-way-latency between machines
- Identify concurrent events with one-way-latency consideration
- Trace chronological order for sequence of events
- Benchmark machines by precise runtime measurement
- Directly utilize RDTSC for maximum precision in event timestamping

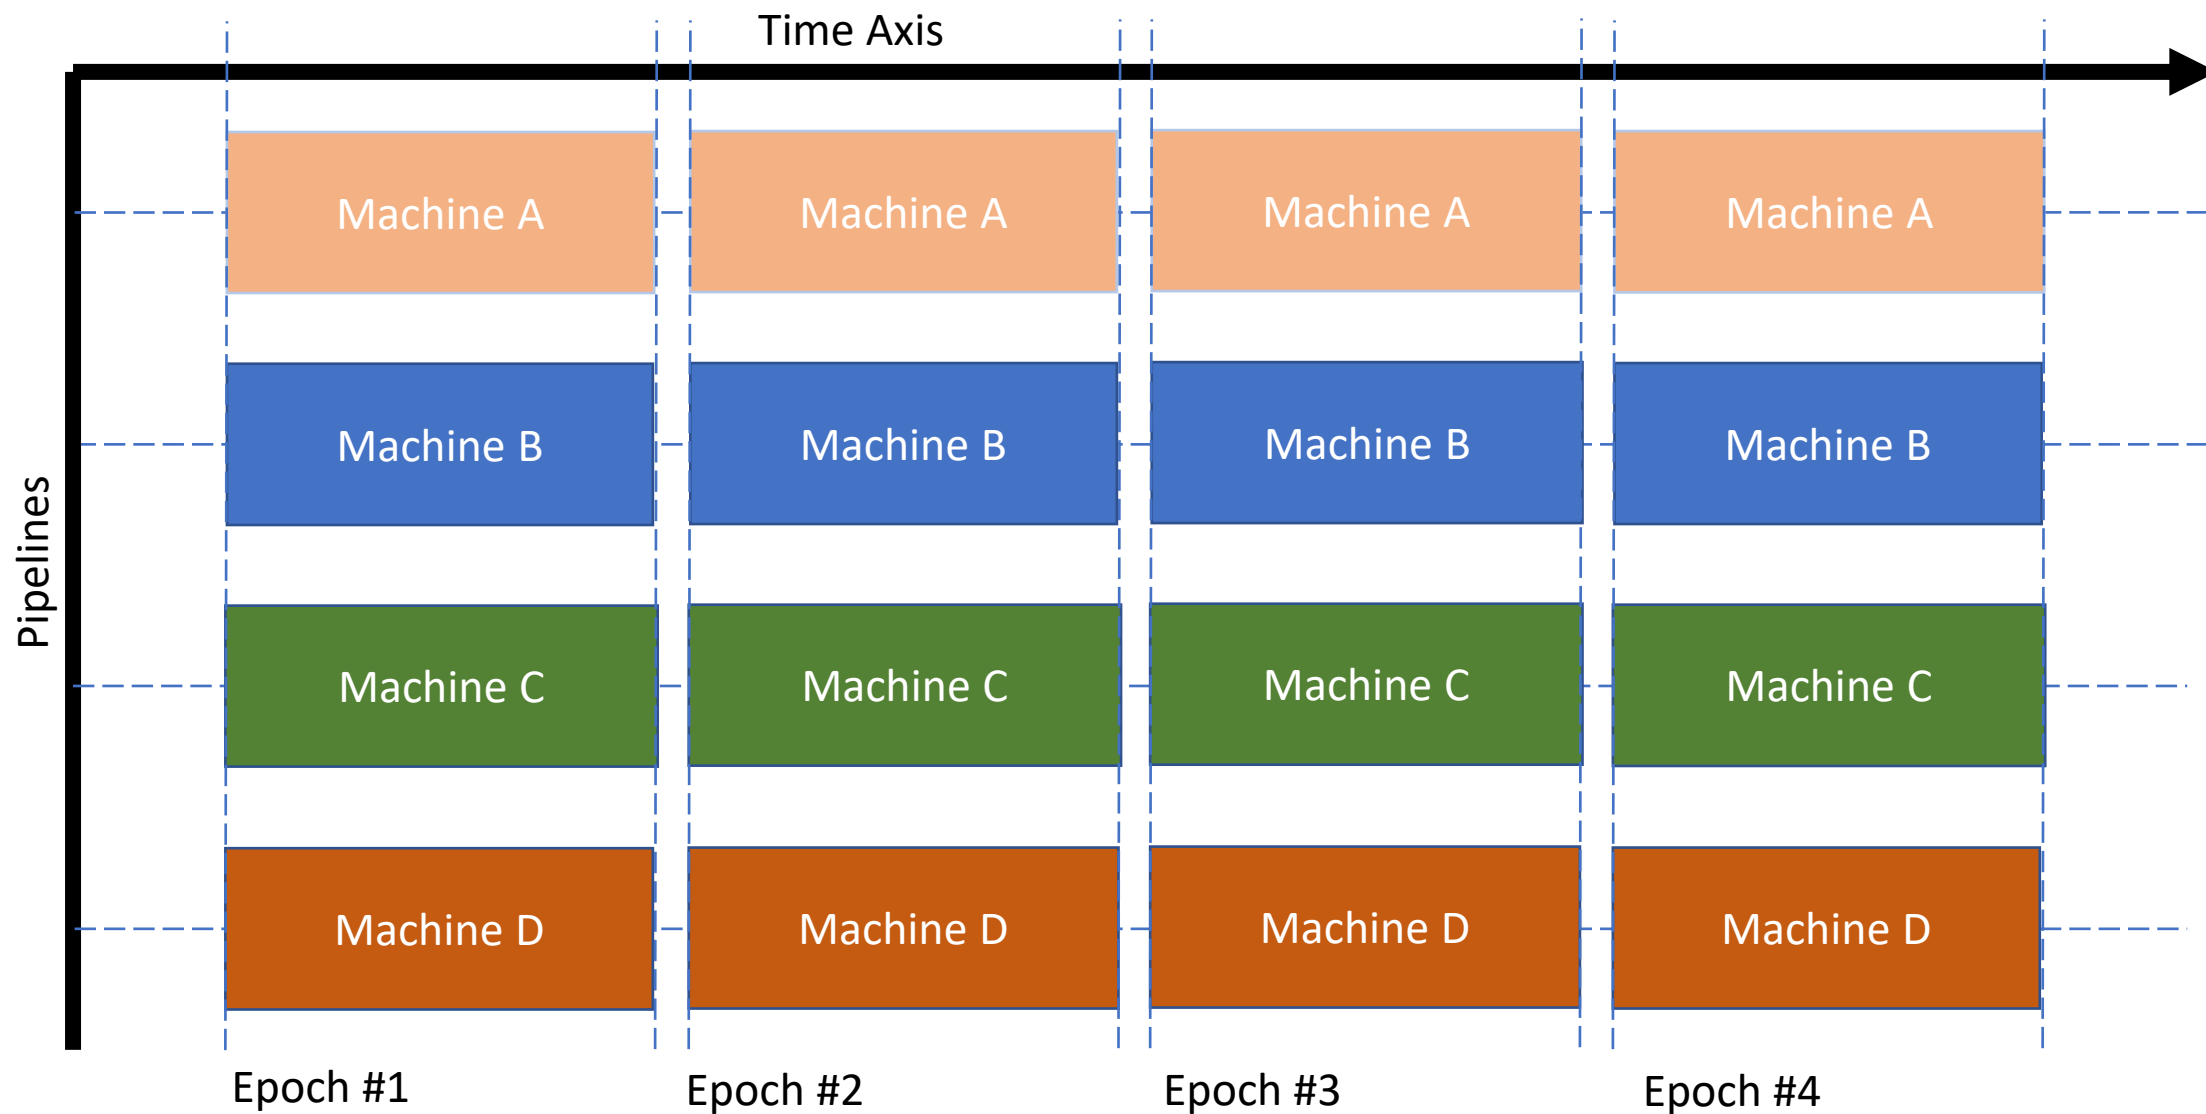
Runtime Difference in a Pipeline




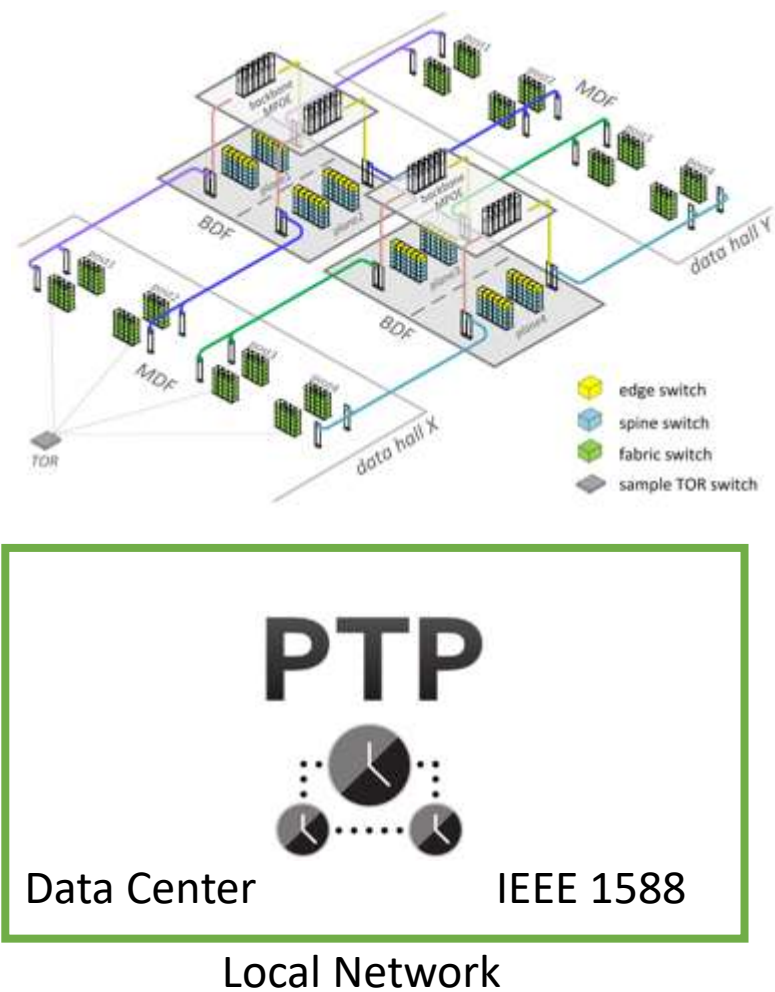

Runtime Difference in a Pipeline



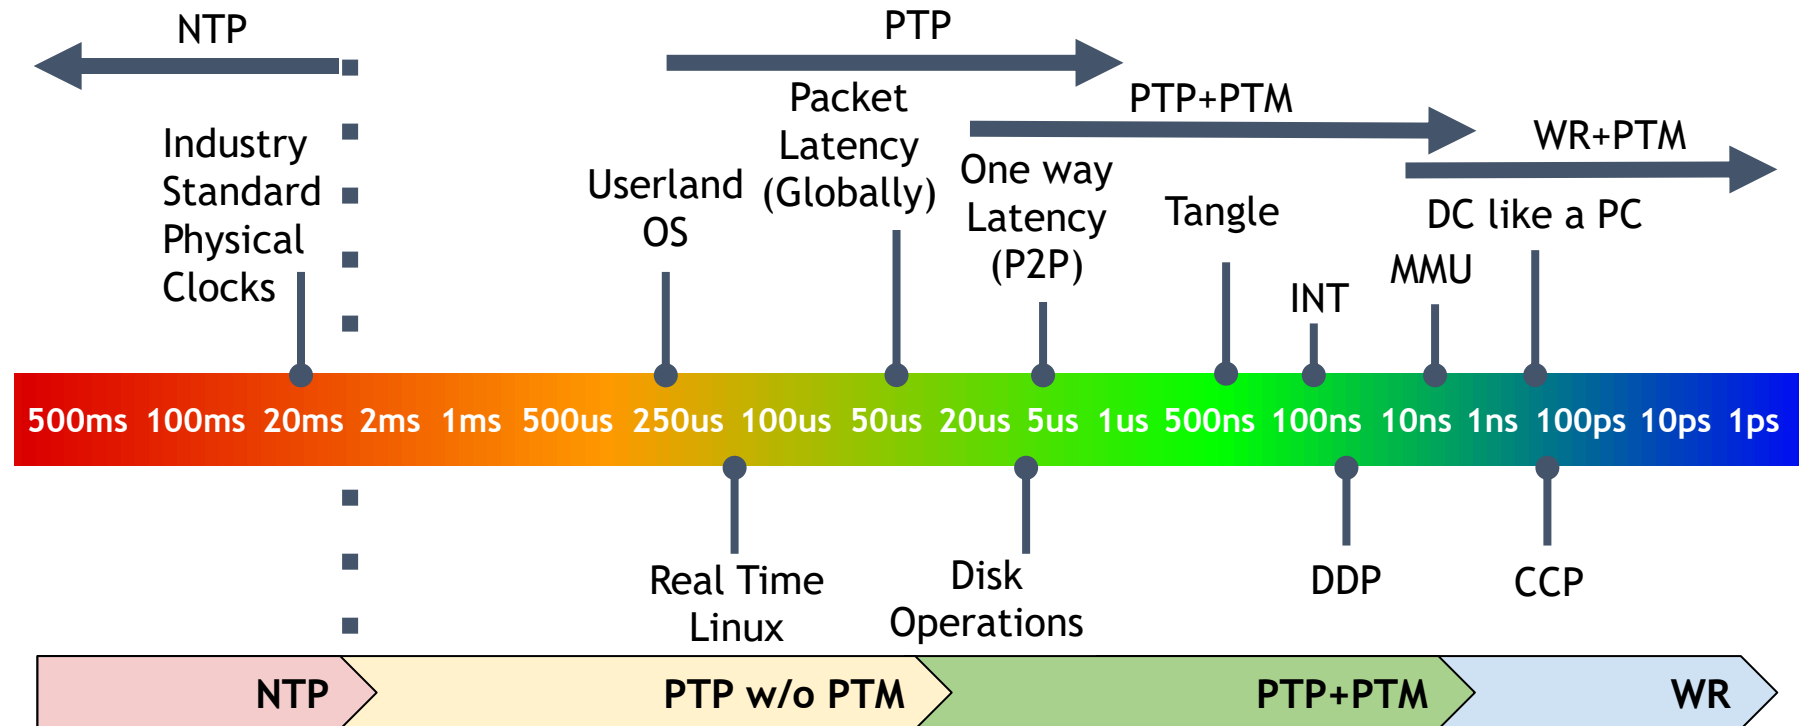
Align and Calibrate Machines in a Pipeline



Precision Time Sync Across Different Domains




Time Precision and Applications




Open Time Server

How to sync a Datacenter?



Traditional

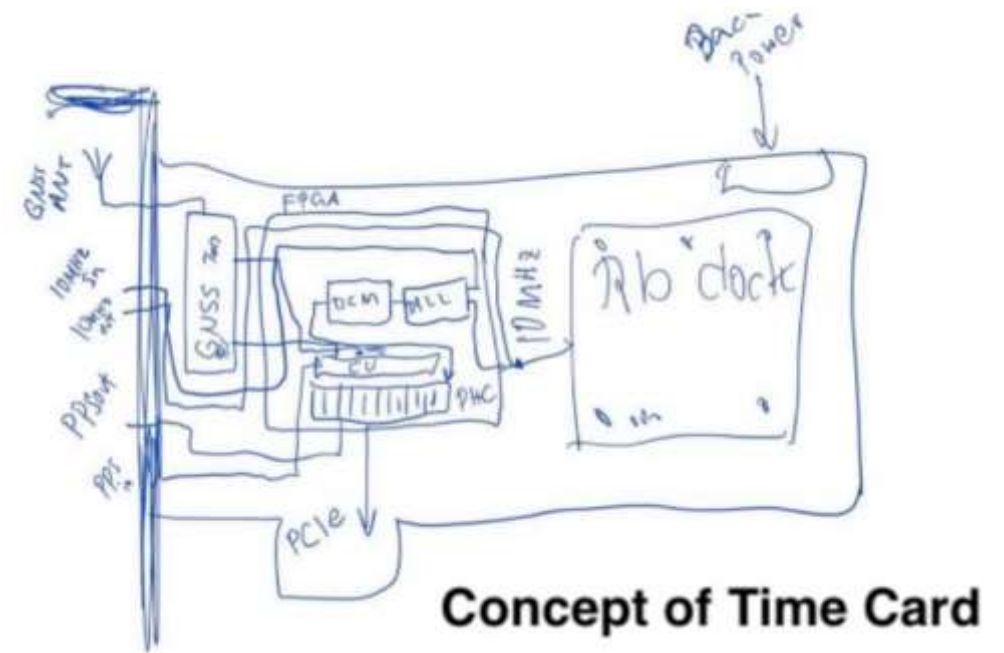


Open Time Server





Time Card


- Concept (2020)
- First Prototype (2021)
- Industry Adoption (2022)



Time Card Family



Latest Time Card




Simplified Objectives for P3335

Produce a guide to improve interoperability, compatibility and comparability between Time Cards manufactured by different vendors.

- Set a structure to define various sub systems
- Set some sub standards for plug-in modules (like GNSS and OSC)
- Define interfaces (hardware: memory addresses, drivers ...)
- Suggest testing methods and comparable benchmark metrics
- Provide blueprints to assist further development of Time Cards
- Do all this work in a structure that maximizes people's engagement

Proposed Approach

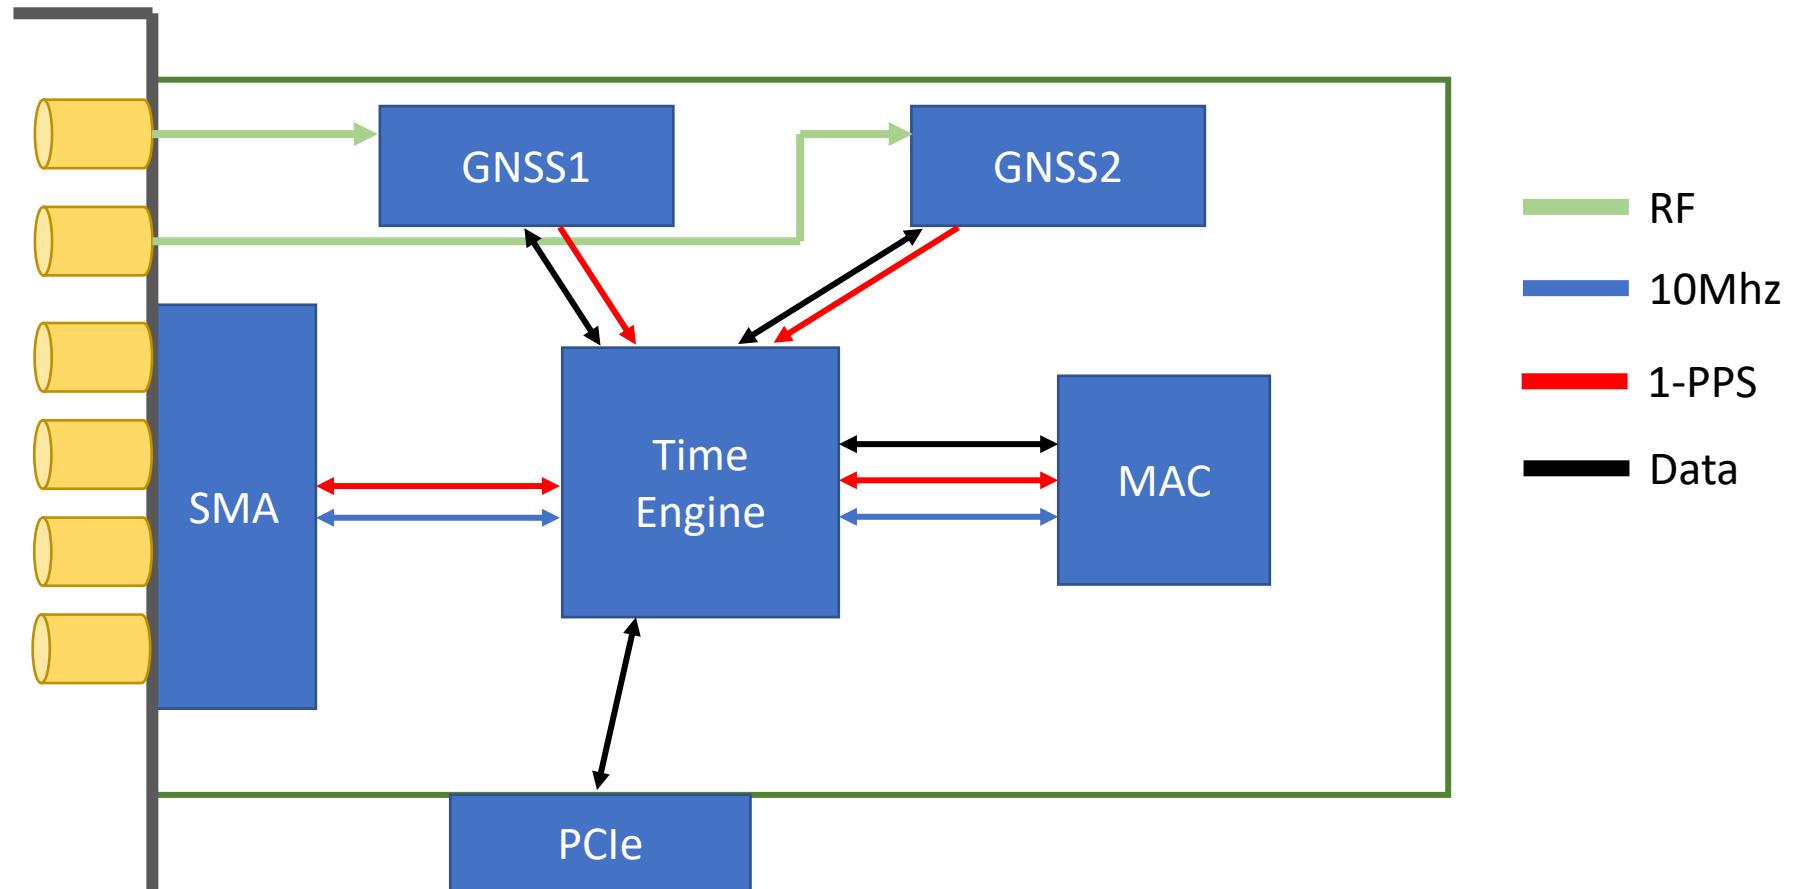
- We already have something that works (various Time Cards)
- Turn the Time Card to a general case
- Study the general case
- Dissect the work in subgroups (Diverge)
- Work in parallel
- Bring things together (Converge)
- Combing through the work (coherence)



Goal: Embrace Diversity, Ensure Interoperability

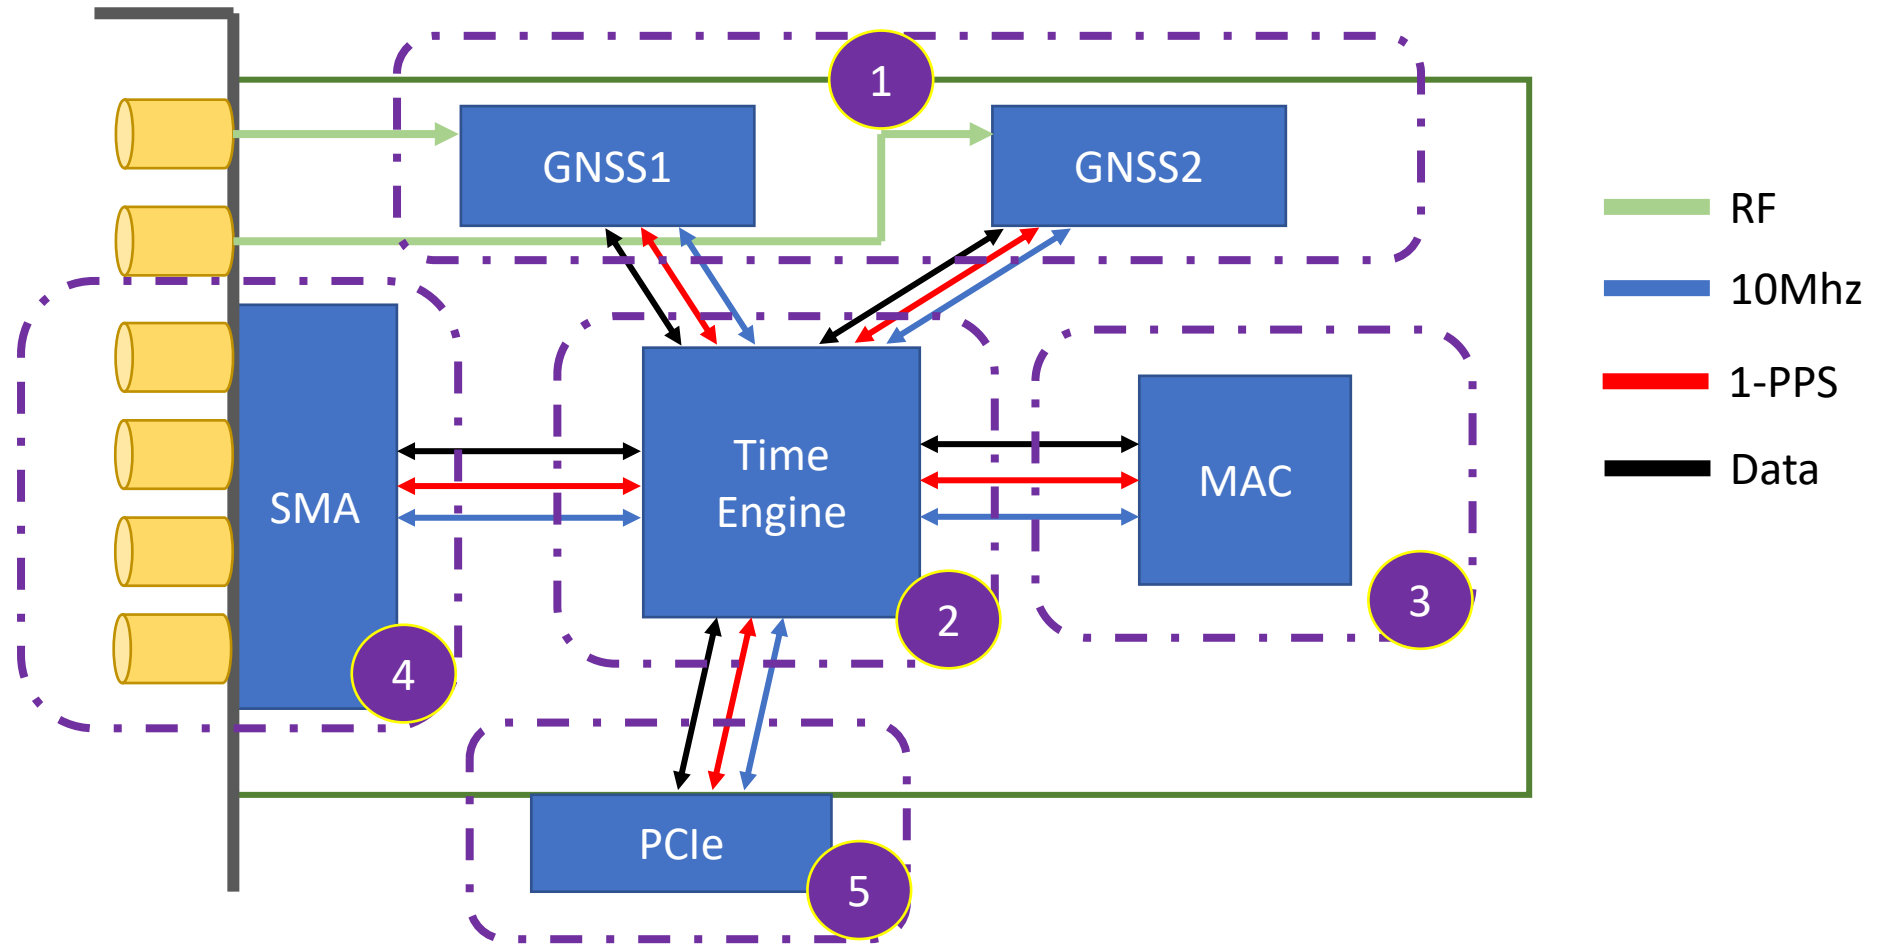
Current Time Card

- Dual GNSS PCIe card with a MAC

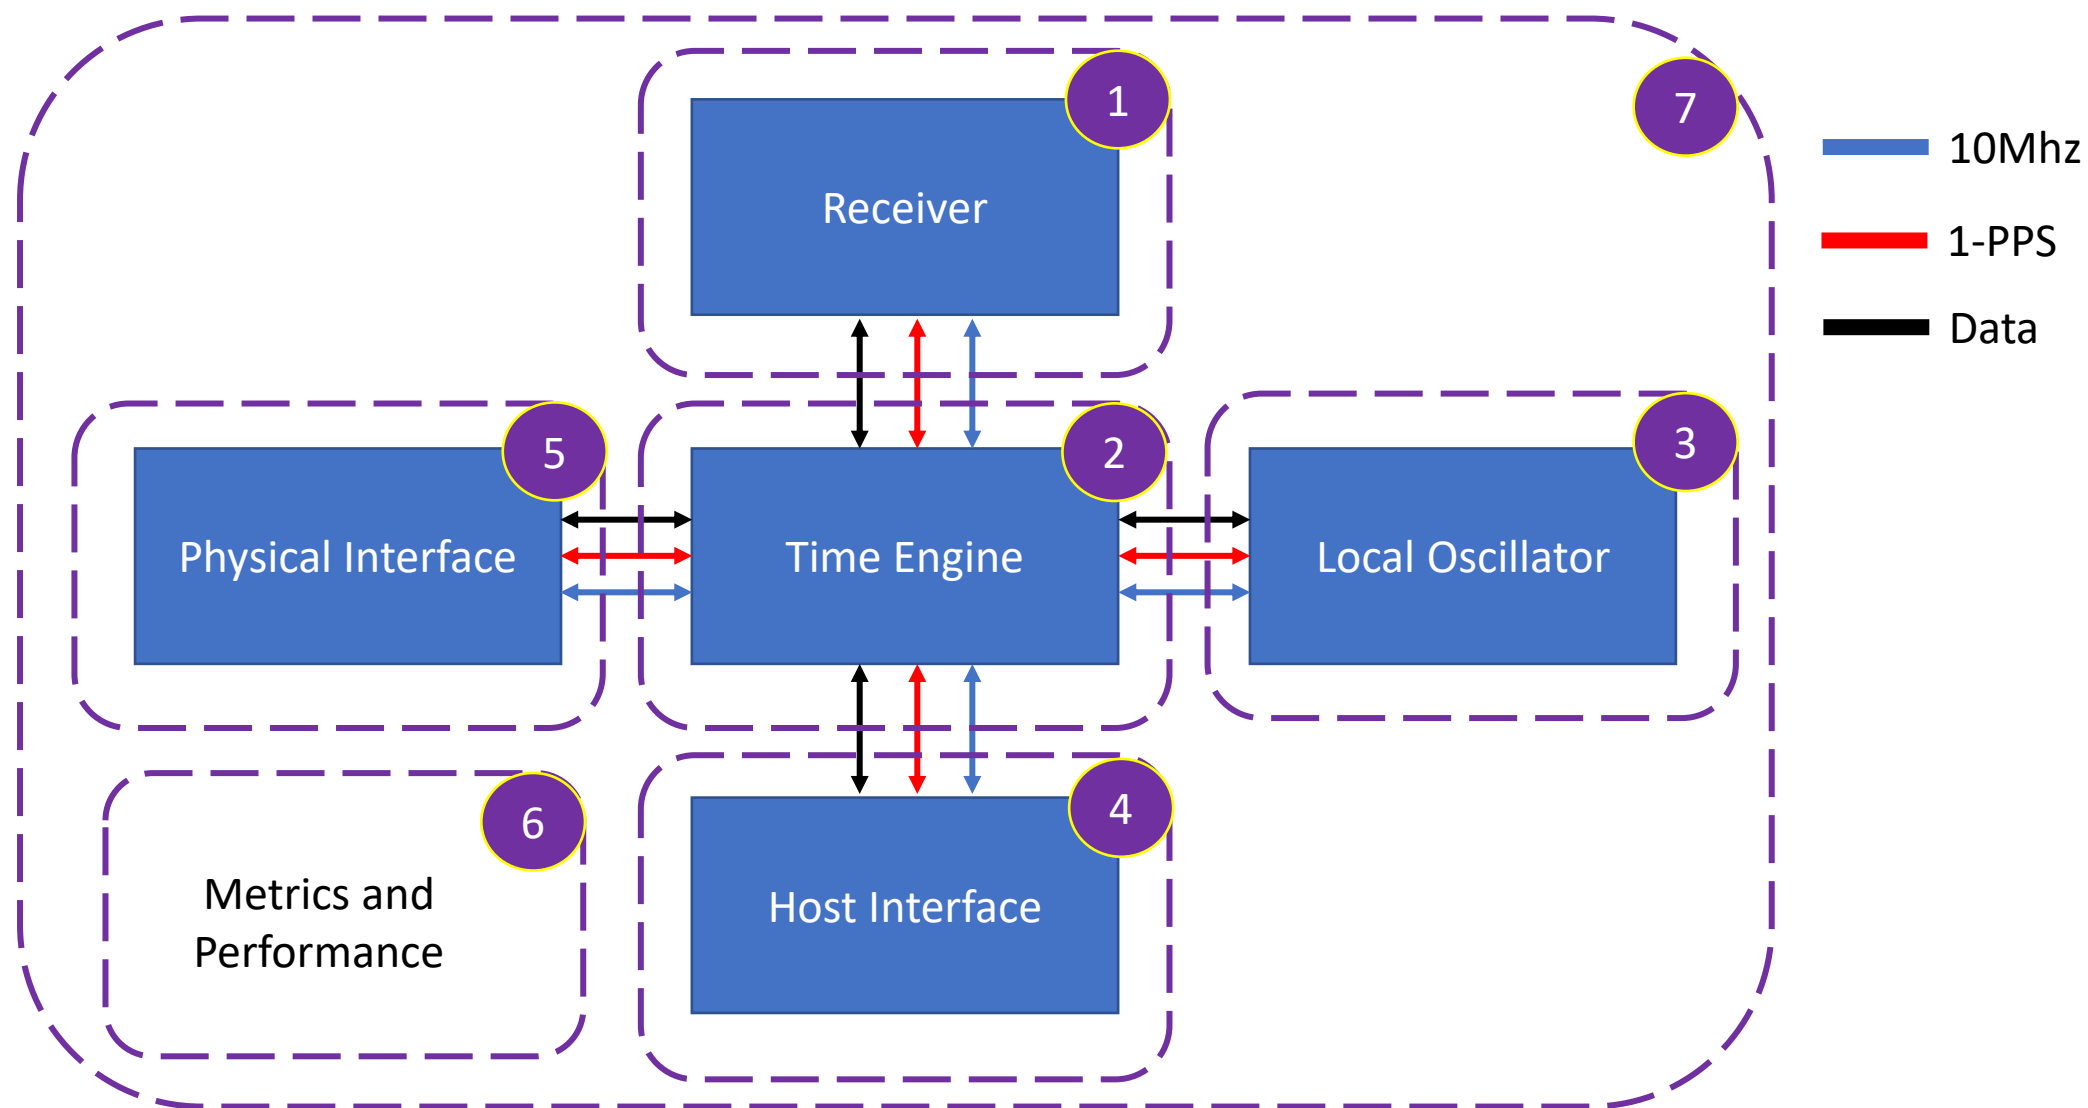


Time Card Concept

- All possible pathways



Focus Areas



Thank you

Find out more on:

www.ocptap.com

www.timecard.ch

www.timingcard.com

www.opentimeserver.com